

LOOM

A cross-portfolio sourcing intelligence layer

30x the buying power. 1x the negotiating leverage.

Built at 11 brands. Designed for 30+.

*“AI will redefine agility for business
— making paths to growth speedy
and seamless.”*

Think9, Principles

Scale that exists on paper, not in the room

15–25% rise in packaging input costs, 2026. Margin defence is a this-quarter problem.

30x

buying power

1x

negotiating leverage

- Each brand sources independently — correct for the brand, ruinous for the portfolio
- Not a discipline problem, and not solved by a shared spreadsheet
- Quick commerce compounds it: every brand now runs q-commerce packs plus D2C plus retail

Because at the moment a brand negotiates, nobody in the building can see what the other brands pay for the same physical object.

Think9 is building toward 30+ consumer brands. Today, four of them buy the same 50ml amber glass bottle — from different vendors, at different prices.

Nobody knows. In the quotes, it has five different names.

THE HARD PROBLEM

One bottle. Five names.

50ml Amber Glass Bottle, 20mm neck

Neude · PDF

Amber Boston Round 50 ML (20/400)

Beauty by Bie · PDF

GB-AMB-50-20N

Panchamrit · WhatsApp

Glass bottle - amber - 50cc - neck 20mm - w/o cap

Neude · email

ਔਂਬਰ ਗਲਾਸ ਬੋਟਲ 50ml 20mm

Goodbug · WhatsApp

ONE PHYSICAL OBJECT

You cannot combine orders for things you cannot prove are the same thing.

So the bottleneck is not price comparison. It is SKU identity resolution — and that is what a script cannot do and a model against a spec graph can.

Nobody wrote these to be inconsistent. They arrived from four vendors across PDF, email and WhatsApp.

A surcharge we pay only because we order apart

“Suppliers may accept 500 units if buyers agree to pay 20–30% more per unit.”

Industry reporting on 2026 packaging procurement

4 ×

Ordering separately — four brands below MOQ, each paying the sub-MOQ premium, four separate times.

0 ×

Ordering together — one consolidated order clears MOQ. The premium does not apply at all.

The saving is not “we negotiated harder.” It is “we stopped paying a surcharge that existed only because we ordered separately.”

A script cannot do it

No fixed schema. ~200 vendor layouts, new ones constantly. Any regex parser is obsolete within a quarter.

A team cannot do it fast enough

Quotes expire in 15–30 days. A human queue is exactly the latency that destroys the value.

It is a loop, not a transform

Low-confidence extraction triggers a clarifying action; resolution decides match-or-create against growing state.

Six layers, three human gates

L0	INGEST	Email, WhatsApp, file drop, ERP export — stored with provenance	
L1	EXTRACTION	Per-field confidence and source-span citation. Never infers a missing value	ROUTING
L2	RESOLUTION	Line item → canonical SKU. Normalization, tolerance rules, fuzzy fallback	GATE 1
L3	SPEC GRAPH	Canonical SKUs, vendors keyed on GSTIN, price history, procurement mode	
L4	OPPORTUNITY	Bundling · price outlier · concentration · lead-time drift · expiry	
L5	BRIEF	The negotiation packet a human takes into the call, reasoning shown	
L6	DECISION	Category manager approves. Outcome writes back — the graph learns	GATES 2 & 3

Confidence routing

≥ 0.90
auto-advance

0.60–0.90
human review

< 0.60
reject

Three gates

- 1 Canonical creation
- 2 Spend recommendation
- 3 Vendor communication

Human effort scales with ambiguity, not volume.

The agent never commits spend and never contacts a vendor.

Value that compounds instead of adding up

Most AI proposals

Linear

Help one team do one thing faster. Deploy to 30 brands, get 30 units of value, pay 30 units of cost.

LOOM

Superlinear

Every quote any brand receives makes the graph denser — which makes every future negotiation across all other brands better informed.

The next brand starts at portfolio-best

- Brand 12 does not negotiate up from scratch — it inherits portfolio-best pricing, qualified vendors and known lead times on day one.
- Think9 grows partly by acquisition. An acquired brand arrives with its own vendor list and its own naming conventions — canonical resolution is exactly the tool that absorbs it.
- This only works if the system sits at the centre. Which is what “centralized intelligence layer” asks for.

It runs. No API key, 1.8 seconds.

12

vendor artifacts

27

line items

23

resolved (85%)

4

held for a human

GLS-AMB-050-20

Five names → one canonical SKU

Neude	Beauty by Bie	Panchamrit	Goodbug
800	600	450	700
Rs 20.75	Rs 21.50	Rs 22.50	Rs 22.80

2,550 units of requirement vs MOQ 2,500 → clears

Rs 13,255

saved on this SKU by consolidating four brands into one order

The refusal matters as much as the match

A 48ml bottle scored 85.3% description similarity against the 50ml canonical — essentially the fuzzy-match threshold. A text matcher merges them. LOOM refuses: volume 48 ≠ 50, both stated. Had it merged, it would have added 1,000 phantom units to the bundle above.

Rs 31,672.60

identified across six audited findings, from 12 quotes

Resolver written before the data existed. Untuned.

Risk findings are excluded from the total — they quantify exposure against spend already counted, not new money.

Thirty days, and a deliberately boring stack



Scoped to packaging — highest cross-brand overlap, no regulatory sign-off to slow the loop. Day 16-23 is the milestone: one real saving, not a dashboard.

What would make me wrong

Wrong canonical merge	Two different SKUs unified, and every recommendation built on a false equivalence.	Gate 1. Reversible merges with audit trail. Precision weighted above recall.
Contract manufacturers buy the packaging	If the co-packer procures, the brand isn't buying it and there is nothing to bundle.	Every brand-SKU edge carries procurement mode. Direct-buy bundles; CM-embedded gets rate-benchmarked.
Vendors learn we bundle	Base quotes inflate in anticipation of consolidated volume.	Do not disclose portfolio volume until commit. Maintain and use BATNA vendors.
Over-consolidation	Single-sourcing a bundled SKU creates a new portfolio-level supply risk.	Concentration is a first-class alert, surfaced alongside every saving.

An architecture that has not enumerated its failure modes has not been designed.